

problem statement

To find the best fit of the dataset

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import preprocessing, svm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

```
In [2]: df=pd.read_csv(r"C:\Users\LENOVO\Downloads\fiat500_VehicleSelection_Dataset (2
df
```

Out[2]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon
0	1	lounge	51	882	25000	1	44.907242	8.611560
1	2	pop	51	1186	32500	1	45.666359	12.241890
2	3	sport	74	4658	142228	1	45.503300	11.417840
3	4	lounge	51	2739	160000	1	40.633171	17.634609
4	5	pop	73	3074	106880	1	41.903221	12.495650
...
1533	1534	sport	51	3712	115280	1	45.069679	7.704920
1534	1535	lounge	74	3835	112000	1	45.845692	8.666870
1535	1536	pop	51	2223	60457	1	45.481541	9.413480
1536	1537	lounge	51	2557	80750	1	45.000702	7.682270
1537	1538	pop	51	1766	54276	1	40.323410	17.568270

1538 rows × 9 columns



```
In [3]: df=df[['age_in_days','km']]
df.columns=['Age','Km']
```


```
In [4]: df.head(10)
```

Out[4]:

	Age	Km
0	882	25000
1	1186	32500
2	4658	142228
3	2739	160000
4	3074	106880
5	3623	70225
6	731	11600
7	1521	49076
8	4049	76000
9	3653	89000

```
In [5]: sns.lmplot(x="Age",y="Km",data=df,order=2,ci=None)
```

Out[5]: <seaborn.axisgrid.FacetGrid at 0x11b30832110>



In [6]: `df.describe()`

Out[6]:

	Age	Km
count	1538.000000	1538.000000
mean	1650.980494	53396.011704
std	1289.522278	40046.830723
min	366.000000	1232.000000
25%	670.000000	20006.250000
50%	1035.000000	39031.000000
75%	2616.000000	79667.750000
max	4658.000000	235000.000000

In [7]: `df.fillna(method='ffill',inplace=True)`

C:\Users\LENOVO\AppData\Local\Temp\ipykernel_6124\3337295870.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

`df.fillna(method='ffill',inplace=True)`

In [8]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1538 entries, 0 to 1537
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -
0   Age      1538 non-null      int64
1   Km       1538 non-null      int64
dtypes: int64(2)
memory usage: 24.2 KB
```

In [9]: `df.fillna(method='ffill',inplace=True)`

C:\Users\LENOVO\AppData\Local\Temp\ipykernel_6124\4116506308.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

`df.fillna(method='ffill',inplace=True)`

```
In [10]: x=np.array(df['Age']).reshape(-1,1)
y=np.array(df['Km']).reshape(-1,1)
```

```
In [11]: df.dropna(inplace=True)
```

C:\Users\LENOVO\AppData\Local\Temp\ipykernel_6124\1379821321.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame


See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df.dropna(inplace=True)
```

```
In [12]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
regr=LinearRegression()
regr.fit(x_train,y_train)
print(regr.score(x_test,y_test))
```


0.6946203593238638

```
In [13]: y_pred=regr.predict(x_test)
plt.scatter(x_test,y_test,color='b')
plt.plot(x_test,y_pred,color='k')
plt.show()
```



```
In [14]: df400=df[:][:400]  
sns.lmplot(x="Age",y="Km",data=df400,order=1,ci=None)
```

Out[14]: <seaborn.axisgrid.FacetGrid at 0x11b40315550>




```
In [15]: x=np.array(df400['Age']).reshape(-1,1)  
y=np.array(df400['Km']).reshape(-1,1)
```

```
In [16]: df400.dropna(inplace=True)  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)  
regr=LinearRegression()  
regr.fit(x_train,y_train)  
print("Regression:",regr.score(x_test,y_test))
```

Regression: 0.7615536976421272

```
In [17]: y_pred=regr.predict(x_test)
plt.scatter(x_test,y_test,color='g')
plt.plot(x_test,y_pred,color='k')
plt.show()
```



```
In [18]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
model=LinearRegression()
model.fit(x_train,y_train)
```

```
Out[18]: ▾ LinearRegression
LinearRegression()
```

```
In [19]: y_pred=model.predict(x_test)
r2=r2_score(y_test,y_pred)
print("r2 score:",r2)
```

r2 score: 0.7615536976421272

conclusion:

This is the best fit dataset since we got accuracy value is 76%.

